

# Medical cost Insurance prediction

```
In [44]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.linear_model import LinearRegression
```

```
In [2]: df = pd.read_csv('medical_cost_insurance.csv')
```

```
In [4]: df.head()
```

```
Out[4]:
```

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520

```
In [5]: df.shape
```

```
Out[5]: (1338, 7)
```

```
In [8]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         1338 non-null   int64
1   sex         1338 non-null   object
2   bmi         1338 non-null   float64
3   children    1338 non-null   int64
4   smoker      1338 non-null   object
5   region      1338 non-null   object
6   charges     1338 non-null   float64
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB
```

```
In [6]: df.describe()
```

```
Out[6]:
```

	age	bmi	children	charges
<b>count</b>	1338.000000	1338.000000	1338.000000	1338.000000
<b>mean</b>	39.207025	30.663397	1.094918	13270.422265
<b>std</b>	14.049960	6.098187	1.205493	12110.011237
<b>min</b>	18.000000	15.960000	0.000000	1121.873900
<b>25%</b>	27.000000	26.296250	0.000000	4740.287150
<b>50%</b>	39.000000	30.400000	1.000000	9382.033000
<b>75%</b>	51.000000	34.693750	2.000000	16639.912515
<b>max</b>	64.000000	53.130000	5.000000	63770.428010

```
In [7]: df.isnull().sum()
```


```
Out[7]: age      0
sex        0
bmi        0
children   0
smoker     0
region     0
charges    0
dtype: int64
```

```
In [10]: plt.figure(figsize=(10, 6))
sns.barplot(x='age', y='charges', data=df, estimator=np.mean, ci=None)
plt.title('Bar Plot of Average Charges by Age Group')
plt.xlabel('Age Group')
plt.ylabel('Average Charges')
plt.show()
```

C:\Users\99Minds-1\AppData\Local\Temp\ipykernel\_13464\2043728585.py:2: FutureWarning:


The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(x='age', y='charges', data=df, estimator=np.mean, ci=None)
```




numerical variable plots


```
In [29]: plt.figure(figsize=(10,5))
plt.hist(df['age'], edgecolor='white', label='d',color='c')
plt.xlabel("Age")
plt.ylabel("Count")
plt.title = ('Age Distrubtion in the Dataset')
```



```
In [30]: plt.figure(figsize=(10,5))
plt.hist(df['bmi'], edgecolor='white', label='d',color='c')
plt.xlabel("BMI")
plt.ylabel("Count")
plt.title = ('BMI Distrubtion in the Dataset')
```




```
In [31]: plt.figure(figsize=(10,5))
plt.hist(df['children'], edgecolor='white', label='d',color='c')
plt.xlabel("Number of Children")
plt.ylabel("Count")
plt.title = ('Children Distrubtion in the Dataset')
```




Categorical variable plots

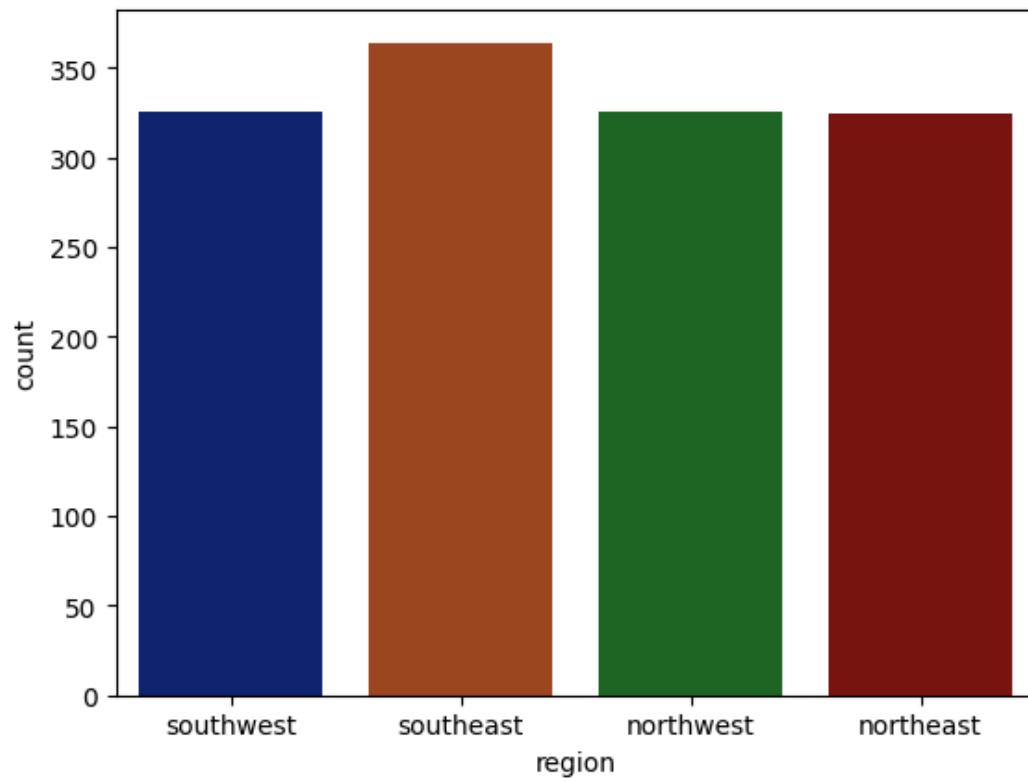
```
In [32]: sex_plot =sns.countplot(x='sex',data=df,palette='dark')
```




```
In [33]: sex_plot = sns.countplot(x='smoker', data=df, palette='dark')
```



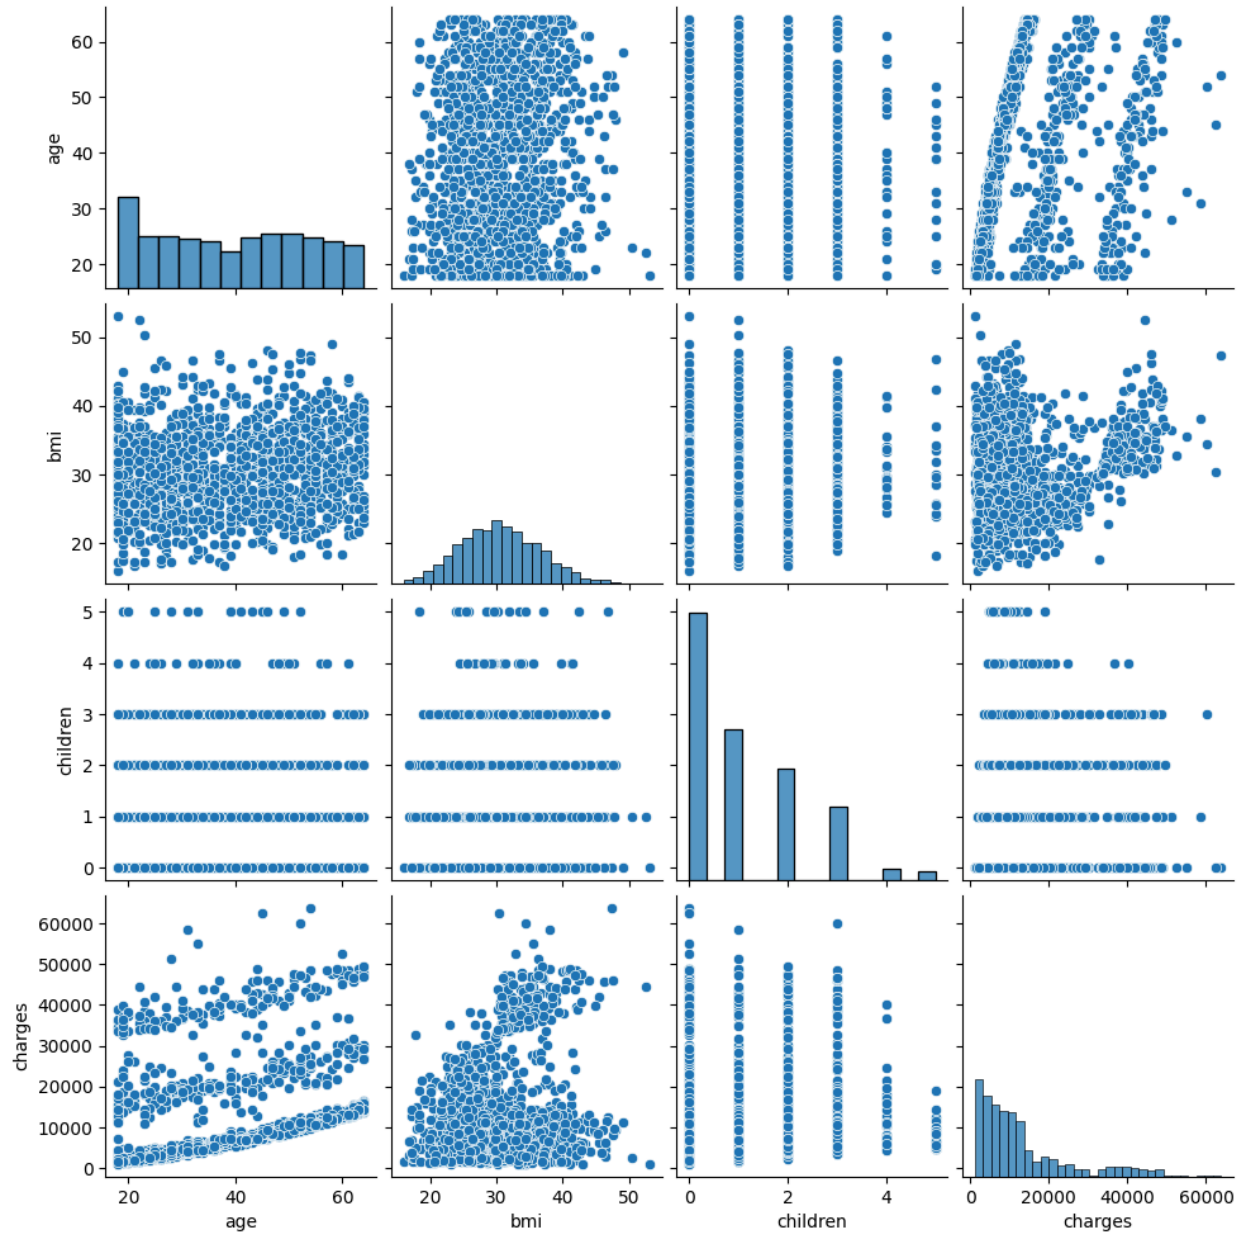
```
In [34]: sex_plot = sns.countplot(x='region', data=df, palette='dark')
```



```
In [12]: plt.figure(figsize=(10, 8))  
sns.heatmap(df[['age', 'bmi', 'children', 'charges']].corr(), annot=True, cmap='coolwarm',  
plt.title('Correlation Matrix')  
plt.show()
```



```
In [13]: sns.pairplot(df)
plt.show()
```



```
In [20]: # Encode categorical variables
df1 = pd.get_dummies(df, dtype=int)
```



In [21]: df1

Out[21]:

	age	bmi	children	charges	sex_female	sex_male	smoker_no	smoker_yes	region_northeast
0	19	27.900	0	16884.92400	1	0	0	1	0
1	18	33.770	1	1725.55230	0	1	1	0	0
2	28	33.000	3	4449.46200	0	1	1	0	0
3	33	22.705	0	21984.47061	0	1	1	0	0
4	32	28.880	0	3866.85520	0	1	1	0	0
...	...	...	...	...	...	...	...	...	...
1333	50	30.970	3	10600.54830	0	1	1	0	0
1334	18	31.920	0	2205.98080	1	0	1	0	1
1335	18	36.850	0	1629.83350	1	0	1	0	0
1336	21	25.800	0	2007.94500	1	0	1	0	0
1337	61	29.070	0	29141.36030	1	0	0	1	0

1338 rows × 12 columns

```
In [25]: # Define features and target
X = df1.drop('charges', axis=1)
y = df1['charges']
```

```
In [26]: # Split the data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [27]: model = LinearRegression()

# Train the model
model.fit(X_train, y_train)
```

```
Out[27]: ▼ LinearRegression
LinearRegression()
```

```
In [38]: from sklearn.metrics import mean_absolute_error, r2_score

# Predict on the test set
y_pred = model.predict(X_test)

# Calculate metrics
mae = mean_absolute_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f'Mean Absolute Error: {mae}')
print(f'R2 Score: {r2}')
```

Mean Absolute Error: 4181.194473753645  
R2 Score: 0.7835929767120723

```
In [37]: from xgboost import XGBRegressor
xgb = XGBRegressor(objective='reg:squarederror', colsample_bytree = 0.3, learning_rate
              max_depth = 5, alpha = 10, n_estimators = 100)
xgb.fit(X_train, y_train)
```

```
Out[37]: XGBRegressor
XGBRegressor(alpha=10, base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=0.3, early_stopping_rounds=None,
              enable_categorical=False, eval_metric=None, feature_types=None,
              gamma=None, gpu_id=None, grow_policy=None, importance_type=None,
              interaction_constraints=None, learning_rate=0.1, max_bin=None,
              max_cat_threshold=None, max_cat_to_onehot=None,
              max_delta_step=None, max_depth=5, max_leaves=None,
              min_child_weight=None, missing=nan, monotone_constraints=None,
              n_estimators=100, n_jobs=None, num_parallel_tree=None,
```

```
In [42]: y_pred = xgb.predict(X_test)
```

```
In [45]: mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
print(f'Mean Squared Error: {mse}')
print(f'R-squared: {r2}')
```

Mean Squared Error: 22785810.99568766  
R-squared: 0.8532302919531741

```
In [47]: plt.scatter(y_test, y_pred)
plt.xlabel('Actual Charges')
plt.ylabel('Predicted Charges')
plt.title('Actual vs Predicted Charges')
plt.show()
```

-----


**TypeError**

Traceback (most recent call last)

Cell In[47], line 4

```
2 plt.xlabel('Actual Charges')
3 plt.ylabel('Predicted Charges')
----> 4 plt.title('Actual vs Predicted Charges')
5 plt.show()
```

**TypeError:** 'str' object is not callable



## Thankyou

In [ ]: